

Specifying Model

- Structure Quadruple $[D, Q, F, R(q_i, d_j)]$
 - D = Representation of documents
 - Q = Representation of queries
 - F = Framework for modeling representations and their relationships
 - Standard language/algebra/impl. type for translation to provide semantics
 - Evaluation w.r.t. “direct” semantics through benchmarks
 - R = Ranking function that associates a real number with a query-doc pair

About index terms

- Each document represented by a set of representative keywords or index terms
 - Index terms meant to capture document's main themes or semantics.
 - Usually, index terms are nouns because nouns have meaning by themselves.
 - However, search engines assume that all words are index terms (full text representation)
- T1 = “conference”
- T2 = “crime”
- Adjectives, adverbs, conjunction, etc not useful.

Notations/Conventions

- K_i is an index term
- d_j is a document
- t is the total number of docs
- $K = (k_1, k_2, \dots, k_t)$ is the set of all index terms
- $w_{ij} \geq 0$ is the weight associated with (k_i, d_j)
 - $w_{ij} = 0$ if the term is not in the doc
- $vec(d_j) = (w_{1j}, w_{2j}, \dots, w_{tj})$ is the weight vector associated with the document d_j
- $gi(vec(d_j)) = w_{ij}$ is the function which returns the weight associated with the pair (k_i, d_j)

The Boolean Model

- Simple model based on set theory
- Queries and documents specified as boolean expressions
 - precise semantics
 - *E.g.*, $q = ka \wedge (kb \vee \neg kc)$
- Terms are either present or absent. Thus, $w_{ij} \in \{0,1\}$

Example


- $q = ka \wedge (kb \vee \neg kc)$
- $vec(qdnf) = (1,1,1) \vee (1,1,0) \vee (1,0,0)$
 - Disjunctive Normal Form
- $vec(qcc) = (1,1,0)$
 - Conjunctive component
- Similar/Matching documents
 - $md1 = [ka \ ka \ d \ e] \Rightarrow (1,0,0)$
 - $md2 = [ka \ kb \ kc] \Rightarrow (1,1,1)$
- Unmatched documents
 - $ud1 = [ka \ kc] \Rightarrow (1,0,1)$
 - $ud2 = [d] \Rightarrow (0,0,0)$

Similarity/Matching function

$$\begin{aligned} \text{sim}(q, dj) &= 1 \text{ if } \text{vec}(dj) \in \text{vec}(qdnf)) \\ &0 \text{ otherwise} \end{aligned}$$

- *Requires coercion for accuracy*

Venn Diagram



$$q = ka \wedge (kb \vee \neg kc)$$

Drawback of Boolean model

- Expressive power of boolean expressions to capture information need and document semantics inadequate
- Retrieval based on binary decision criteria (with no partial match) does not reflect our intuitions behind relevance adequately
- As a result
 - Answer set contains either too few or too many documents in response to a user query
 - No ranking of documents

Vector Model


- Task:
 - Document collection
 - Query specifies information need: free text
 - Relevance judgments: depends upon the weighting scheme for all docs
- Word evidence: Bag of words
 - No ordering information

Vector Space Model

- Represent documents and queries as
 - Vectors of term-based features
 - Features: tied to occurrence of terms in collection
 - E.g. $\vec{d}_j = (t_{1,j}, t_{2,j}, \dots, t_{N,j}); \vec{q}_k = (t_{1,k}, t_{2,k}, \dots, t_{N,k})$
- Solution 1: Binary features: $t=1$ if presence, 0 otherwise
 - Similarity: number of terms in common
 - Dot product

$$\text{sim}(\vec{q}_k, \vec{d}_j) = \sum_{i=1}^N t_{i,k} t_{i,j}$$

The Vector Model: Example I



	k1	k2	k3		$q \bullet d_j$
d1	1	0	1		2
d2	1	0	0		1
d3	0	1	1		2
d4	1	0	0		1
d5	1	1	1		3
d6	1	1	0		2
d7	0	1	0		1
q	1	1	1		

Vector Space Model II


- Problem: Not all terms equally interesting
 - E.g. “accuracy” vs “crime”

$$\vec{d}_j = (w_{1,j}, w_{2,j}, \dots, w_{N,j}); \vec{q}_k = (w_{1,k}, w_{2,k}, \dots, w_{N,k})$$

- Solution: Replace binary term features with weights
 - Document collection: term-by-document matrix
 - View as vector in multidimensional space
 - Nearby vectors are related
 - Normalize for vector length

Cosine similarity

- Distance between vectors d_1 and d_2 captured by the cosine of the angle θ between them.



Queries in the vector space model

Central idea: the query as a vector:

- We regard the query as short document
 - Note that d_q is very sparse!
- We return the documents ranked by the closeness of their vectors to the query, also represented as a vector.

$$\text{sim}(d_j, d_q) = \frac{\vec{d}_j \cdot \vec{d}_q}{\|\vec{d}_j\| \|\vec{d}_q\|} = \frac{\sum_{i=1}^n w_{i,j} w_{i,q}}{\sqrt{\sum_{i=1}^n w_{i,j}^2} \sqrt{\sum_{i=1}^n w_{i,q}^2}}$$

Vector Similarity Computation

- Similarity = Dot product

$$\text{sim}(\vec{q}_k, \vec{d}_j) = \vec{q}_k \bullet \vec{d}_j = \sum_{i=1}^N w_{i,k} w_{i,j}$$

- Normalization:

- Normalize weights in advance
- Normalize post-hoc

$$\text{sim}(\vec{q}_k, \vec{d}_j) = \frac{\sum_{i=1}^N w_{i,k} w_{i,j}}{\sqrt{\sum_{i=1}^N w_{i,k}^2} \sqrt{\sum_{i=1}^N w_{i,j}^2}}$$

- Cosine of angle between two vectors
- The denominator involves the lengths of the vectors.

Computation of weights w_{ij} and w_{iq}

- How to compute the weights w_{ij} and w_{iq} ?
 - quantification of intra-document content (similarity/semantic emphasis)
 - *tf* factor, the *term frequency* within a document
 - quantification of inter-document separation (dis-similarity/significant discriminant)
 - *idf* factor, the *inverse document frequency*
 - $w_{ij} = tf(i,j) * idf(i)$

Weighting scheme

- Let,
 - N be the total number of docs in the collection
 - n_i be the number of docs which contain k_i
 - $freq(i,j)$ raw frequency of k_i within d_j
- A normalized tf factor is given by
 - $f(i,j) = freq(i,j) / \max(freq(l,j))$
 - where the maximum is computed over all terms which occur within the document d_j
- The idf factor is computed as
 - $idf(i) = \log (N/n_i)$
 - the \log makes the values of tf and idf comparable.


Rules:

- WARNING: In a lot of IR literature, “frequency” is used to mean “count”
 - Thus *term frequency* in IR literature is used to mean *number of occurrences* in a doc
 - Not divided by document length (which would actually make it a frequency)

Best weighting scheme


- The best term-weighting schemes use weights which are given by
 - $w_{ij} = f(i,j) * \log(N/n_i)$
 - the strategy is called a *tf-idf* weighting scheme
- For the query term weights, use
 - $w_{iq} = (0.5 + [0.5 * \text{freq}(i,q) / \max(\text{freq}(l,q))]) * \log(N/n_i)$
- The vector model with *tf-idf* weights is a good ranking strategy for general collections.
 - It is also simple and fast to compute.

The Vector Model: Example II



	k1	k2	k3	$q \bullet d_j$
d1	1	0	1	4
d2	1	0	0	1
d3	0	1	1	5
d4	1	0	0	1
d5	1	1	1	6
d6	1	1	0	3
d7	0	1	0	2
q	1	2	3	

The Vector Model: Example III



	k1	k2	k3	$q \bullet d_j$
d1	2	0	1	5
d2	1	0	0	1
d3	0	1	3	11
d4	2	0	0	2
d5	1	2	4	17
d6	1	2	0	5
d7	0	5	0	10
q	1	2	3	

Example 2(Boolean model)

- Consider these documents:
 - **Doc 1** breakthrough drug for schizophrenia
 - **Doc 2** new schizophrenia drug
 - **Doc 3** new approach for treatment of schizophrenia
 - **Doc 4** new hopes for schizophrenia patients
-
- For the document collection, Use and depict the Boolean model and what are the Returned results for these queries:
 - a. schizophrenia AND drug

Example (vector model)

- Q : “gold silver truck”
 - D1 : “shipment of gold damaged in a fire”
 - D2 : “delivery of silver arrived in a silver truck”
 - D3 : “Shipment of gold in a truck”
-
- Find the ranking of the document using vector space model.

Example 2

**Q 1: “About modi interview in politics”
(5)**

D1 : “In the biggest *interview* of 2014 Arnab asks all the questions that India wanted answers from the Gandhi”

D2 : “Interview with BJP leader Narendra Modi | India Insight – Reuters”

D3 : “among all politicians *Modi* is the most polarizing *politician* in India”